

California Law Chatbot - Comprehensive Technical Guide

Updated: November 3, 2025

Table of Contents

1. [Executive Summary](#)
 2. [System Architecture](#)
 3. [Core Components](#)
 4. [CEB RAG Integration](#)
 5. [Data Processing Pipeline](#)
 6. [Multi-Modal Response System](#)
 7. [Anti-Hallucination & Verification](#)
 8. [Multi-Turn Conversation Handling](#)
 9. [API Integrations](#)
 10. [Deployment & Environment](#)
 11. [User Experience Features](#)
 12. [Technical Implementation](#)
 13. [Performance & Cost Analysis](#)
 14. [Limitations & Compliance](#)
-

Executive Summary

The **California Law Chatbot** is a production-ready, enterprise-grade legal research assistant that combines advanced AI models with authoritative legal data sources to provide accurate, verified California law information. This system implements a sophisticated **Generator-Verifier** architecture with multiple layers of validation to minimize AI hallucinations while maintaining conversational fluency and real-time data access.

Key Features

- **CEB RAG Integration:** 77,406 vector embeddings from 2,554 CEB documents across 5 legal verticals (Trusts & Estates, Family Law, Business Litigation, Business Entities, Business Transactions)
- **Multi-Modal Response System:** 3 source modes - CEB Only, AI Only, Hybrid (recommended)
- **Advanced AI Models:** Gemini 2.5 Pro (generator) + Claude Sonnet 4.5 (verifier)
- **Real-Time Legal Sources:** CourtListener (case law), OpenStates/LegiScan (legislation), Google Search grounding
- **Multi-Turn Context:** Intelligent conversation memory with query expansion
- **Comprehensive Verification:** Two-pass claim verification with dynamic confidence gating
- **Vercel Deployment:** Serverless architecture with Upstash Vector database
- **Total Coverage:** 5 CEB verticals + 1M+ CourtListener cases + real-time legislative data

System Maturity

- **Production Ready:** All core features fully implemented and tested
 - **Cost Optimized:** \$0.02-0.30 per comprehensive query (including embeddings)
 - **Scale:** Handles 77,406+ vectors with <2s retrieval latency
 - **Compliance:** California State Bar compliant with mandatory disclaimers
-

System Architecture

High-Level Overview

The system architecture is designed to handle complex legal queries by integrating multiple components:

- React UI:** The front-end interface for users.
- AI Engine:** Handles AI-related tasks, including multi-turn reasoning and context handling.
- Legal APIs:** External sources for legal information.
- Source Mode Selector:** Determines the source type (CEB/AI/Hybrid).
- Handler:** Manages the flow of data between components.
- Multi-Turn:** Handles iterative AI interactions.
- Context:** Manages contextual information across turns.
- Sources:** External legal databases and APIs.
- External Sources:** CEB Vector (CEB RAG + Legislation).
- Database:** Stores historical data and results.
- Generator:** Generates responses using Gemini 2.5 Pro or Claude Sonnet.
- Verifier:** Verifies the generated responses using Google Search.
- Guardrails:** Ensures responses are safe and compliant.
- Gating:** Controls the flow of data through the system.
- Confidence:** Measures the confidence of generated responses.
- Validation:** Checks the validity of generated responses.
- Thresholds:** Sets thresholds for various system parameters.

Core Design Principles

1. Source Authority Hierarchy

- **Level 1:** CEB Practice Guides (authoritative, no verification needed)
 - **Level 2:** Full bill text (OpenStates/Legiscan) - 30% verification threshold
 - **Level 3:** CourtListener case law - 60% verification threshold
 - **Level 4:** Google Search grounding - 20% verification threshold
 - **Level 5:** AI training data - 100% verification required

2. Multi-Modal Processing

- **CEB Only:** Fastest mode, authoritative CEB practice guides only
 - **AI Only:** External APIs + Google Search (full verification)
 - **Hybrid:** CEB-first with AI supplementation (recommended)

3. Context Preservation

- Last 10 messages maintained for conversation memory
 - Intelligent query expansion for vague follow-ups
 - Session-aware source routing and verification

Core Components

1. Frontend (React 19 + TypeScript)

Key Files

Source Mode Selector

Users can switch between three modes:

Mode	Description	Use Case
CEB Only	Authoritative CEB practice guides only	Trusts & Estates, Family Law, Business Litigation questions
Hybrid	CEB + AI sources (recommended)	Comprehensive research combining practice guides + case law
AI Only	External APIs only (no CEB)	General legal research, non-CEB topics

Response Badges

- **CEB VERIFIED** (琥珀色): Authoritative CEB source, no verification needed
- **CourtListener Enhanced** (蓝色): Case law sources included
- **Google Search Grounding** (绿色): Real-time web data used
- **Verification Status** (深灰色): Claim verification coverage

2. Backend Services (TypeScript + Vercel)

Chat Orchestration (`gemini/chatService.ts`)

```
// Main processing pipeline
async sendMessage(message: string, conversationHistory, sourceMode, signature)
  // 1. Context expansion for multi-turn queries
  const expandedQuery = this.expandQueryWithContext(message, conversationHistory)

  // 2. Route to appropriate mode
  switch (sourceMode) {
    case 'ceb-only': return await this.processCEBOnly(expandedQuery, history);
    case 'ai-only': return await this.processAIOnly(expandedQuery, history);
    case 'hybrid': return await this.processHybrid(expandedQuery, history);
  }
}
```

CEB Integration (`api/ceb-search.ts`)

- **Vector Database:** Upstash Vector (serverless, Vercel-compatible)
- **Embedding Model:** OpenAI `text-embedding-3-small` (1536 dimensions, cosine similarity)
- **Namespaces:** 5 separate namespaces per legal vertical
- **Retrieval:** Top-K semantic search with confidence filtering (≥ 0.7 threshold)

AI Pipeline

1. Generator: Gemini 2.5 Pro with Google Search grounding

- Temperature: 0.2 (legal accuracy)
- System prompt: California law expert with mandatory grounding
- Context window: Last 10 messages (multi-turn support)

2. Verifier: Claude Sonnet 4.5

- Claim extraction from generated answer
- Source validation against retrieved documents
- Dynamic verification thresholds

- Response rewriting for unsupported claims

3. Guardrails: Citation validation, legal entity checking

- Validates [1], [2] references match actual sources
 - Flags hallucinated codes/cases
 - Ensures California jurisdiction focus
-

CEB RAG Integration

Architecture Overview

CEB Documents (PDFs) → Processing Pipeline → Vector Database → Response Generation

The architecture diagram illustrates the flow from raw CEB documents to final responses. It starts with a large input of "Source PDFs" (2,554 files), which are processed by "Scripts" (Python) to extract text. This text is then categorized into "5 Legal Verticals": Trusts, Family, Business, etc. The processed text is stored in a "Vector Database". Finally, a "Semantic Search" (API Endpoint) retrieves relevant chunks from the database to generate a "Response" using "Gemini 2.5 Pro". The entire process is managed by "Response Generation" components.

```

graph TD
    A[Source PDFs (2,554 files)] --> B[Scripts (Python)]
    B --> C[5 Legal Verticals: Trusts, Family, Business, etc.]
    C --> D[Vector Database]
    D --> E[Semantic Search (API Endpoint)]
    E --> F[Response Generation (Gemini 2.5 Pro)]
    F --> G[Response]
    
```

Processing Pipeline Details

1. PDF Processing (`scripts/process_ceb_pdfs.py`)

Input: Raw PDF documents from CEB website **Output:** JSONL chunks with metadata

Processing Steps: 1. **Text Extraction:** PyPDF2 + PDFMiner for robust text extraction

Chunking: 1000-token chunks (~4000 characters) with overlap

Metadata Generation: - Source file name - Page number - Section titles - Token count - CEB citation format

Quality Control: - Remove empty pages - Filter low-content chunks (<200 tokens) - Validate PDF parsing errors

Vertical Statistics:

Vertical	PDFs	Chunks	Tokens
Trusts & Estates	1,687	40,263	40.2M
Family Law	243	7,511	7.5M
Business Litigation	323	13,711	13.7M
Business Entities	270	10,766	10.8M
Business Transactions	246	7,517	7.5M
Total	2,554	77,406	77.4M

2. Embedding Generation (`scripts/generate_embeddings.py`)

Model: OpenAI `text-embedding-3-small` **Dimensions:** 1536 (optimized for legal text)
Cost: \$0.00002 per 1K tokens **Total Cost:** ~\$1.55 for initial 77M token embedding

Batch Processing: - 500 chunks per batch (rate limit handling) - Automatic retry on 429 errors - Progress tracking with ETA - Memory optimized for M4 Max (128GB RAM)

3. Vector Upload (`scripts/upload_to_upstash.py`)

Database: Upstash Vector (serverless, 99.99% uptime) **Format:** `{id, vector, metadata, namespace}` **Namespace Strategy:** `ceb_trusts_estates`, `ceb_family_law`, etc. **Metadata Storage:** - Full text (10KB max per chunk) - CEB citation - Page/section numbers - Source file reference - Confidence score

Upload Statistics: - **Total Vectors:** 77,406 - **Upload Time:** 28 minutes (parallelized) - **Success Rate:** 100% (0 failures) - **Storage:** ~2.1GB vector data + metadata

4. Semantic Search Implementation (`api/ceb-search.ts`)

Query Flow:

```
// 1. Query expansion (multi-turn context)
const expandedQuery = expandQueryWithContext(message, conversationHistory);

// 2. Category detection
const category = detectCEBCategory(expandedQuery);

// 3. Embedding generation (OpenAI)
const queryEmbedding = await generateEmbedding(expandedQuery);

// 4. Vector search (Upstash)
const results = await index.query({
  vector: queryEmbedding,
  topK: 5,
  namespace: `ceb_${category}`,
  includeMetadata: true
});

// 5. Confidence filtering (≥ 0.7 threshold)
const highConfidence = results.filter(r => r.metadata.confidence >= 0.7);
```

Category Detection Algorithm: - Keyword matching across 5 verticals - Score-based routing (highest matching vertical) - Default: Trusts & Estates (broadest coverage) - Multi-vertical fallback for ambiguous queries

Multi-Modal Response System

Source Mode Architecture

1. CEB Only Mode (Fastest, Authoritative)

Use Case: Users want answers from CEB practice guides only **Verification:** None required (CEB is authoritative) **Latency:** ~2-4 seconds **Sources:** Upstash Vector search only

Processing Flow:

User Query â†’ Category Detection â†’ CEB Vector Search (topK=5)
â†’ Confidence Filter (≥ 0.7) â†’ Gemini 2.5 Pro Synthesis
â†’ CEB Verified Response (No Verification)

Response Characteristics: - "CEB VERIFIED" badge - Citations: Cal. Prac. Guide: Family Law Â§ 3:45 - Confidence: â‰¥70% similarity threshold - Fallback: "No relevant CEB guidance found"

2. AI Only Mode (Comprehensive Search)

Use Case: General legal research without CEB constraints **Verification:** Full two-pass verification required **Latency:** ~15-25 seconds **Sources:** CourtListener + Legislation APIs + Google Search

Processing Flow:

User Query → Multi-Source Search (CourtListener, OpenStates, Legiscan)
→ Source Pruning (top-K, dedupe, rerank) → Gemini 2.5 Pro Generation
→ Claude Sonnet 4.5 Verification → Confidence Gating
→ Guardrails → Final Response

Verification Threshold: 60% (standard) **Response Characteristics:** - Full verification badges (â€œ/âš) - External source citations - Detailed verification reports

3. Hybrid Mode (Recommended)

Use Case: Comprehensive research combining authoritative + current sources **Verification:** Conditional (CEB bypasses, AI requires) **Latency:** ~10-20 seconds **Sources:** CEB + CourtListener + Legislation + Google Search

Processing Flow:

User Query → Parallel Search (CEB + AI Sources)
→ CEB Priority Ranking → Combined Context Building
→ Gemini 2.5 Pro Synthesis (CEB-first) → Conditional Verification
→ Source Attribution → Final Hybrid Response

Smart Integration Logic:

```

// Priority: CEB > Full Bill Text > Case Law > Google Search
const highConfidenceCEB = cebSources.filter(s => s.confidence >= 0.7);
const needsVerification = aiSources.length > 0 && highConfidenceCEB.length > 0;

if (highConfidenceCEB.length > 0) {
    // CEB dominates - minimal verification
    verificationStatus = 'not_needed';
} else {
    // AI-heavy - full verification
    verificationStatus = await verifyAIResponse();
}

```

Response Characteristics: - CEB sources displayed first with "CEB VERIFIED" badges - AI sources integrated as supplementary context - Mixed verification status based on source composition - Intelligent citation blending

Anti-Hallucination & Verification

Dynamic Confidence Gating

Verification Thresholds

Source Type	Threshold	Rationale
CEB Practice Guides	0%	Authoritative primary sources
Full Bill Text	30%	Official legislative text
CourtListener Cases	60%	Judicial opinions with precedent
Google Search Results	20%	Real-time data with grounding
AI Training Data	100%	Requires full verification

Two-Pass Verification Process

Pass 1: Claim Extraction (Gemini 2.5 Pro)

```

// Extract specific, verifiable claims from generated answer
const claims = VerifierService.extractClaimsFromAnswer(response.text, sources)
// Result: ["Family Code § 4320 lists 14 factors", "Support is tax-deductible"]

```

Pass 2: Source Validation (Claude Sonnet 4.5)

```

// For each claim, validate against available sources
const verificationResults = await verifier.verifyClaims(claims, sources)
// Result: { coverage: 0.85, verified: 12/14, unverified: 2/14 }

```

Confidence Calculation: - **Coverage:** Verified claims ÷ Total claims - **Min Support:** Minimum source references per verified claim - **Ambiguity:** Conflicting information between sources

Guardrail System

1. Citation Validation

```

// Ensure [1], [2] references point to actual sources
const citations = extractCitations(response);
for (const citation of citations) {
  if (citation.index > sources.length) {
    return "BLOCKED: Invalid citation reference";
  }
}

```

2. Jurisdiction Guardrail

```

// Flag non-California sources
if (response.includes("U.S. Supreme Court") && !isFederalRelevant(message))
  return "CAUTION: This response includes federal law. California law may differ";
}

```

3. Temporal Accuracy

```

// Warn about outdated information
const billPattern = /AB|SB|Assembly Bill|Senate Bill/;
if (billPattern.test(message) && !hasRecentData(response)) {
  return "WARNING: Verify current bill status - this may not reflect law";
}

```

Multi-Turn Conversation Handling

Context Expansion Algorithm

Query Expansion Examples:

User Sequence	Original Query	Expanded Query
"What is Penal Code 459?"	"What about 460?"	"What is Penal Code 460?"
"What about 460?"		
"Explain burglary laws"		
"Does it apply to houses?"	"Does it apply to houses?"	"Regarding burglary laws, does it apply to houses?"
"What is Family Code 4320?"		
"What if the marriage was short?"	"What if the marriage was short?"	"What if the marriage was short? (in the context of Family Code 4320)"

Pattern Matching: 1. **Code Section Follow-ups:** Detects code type from previous query, applies to new section number 2. **"Does it/Is it/Can it" Questions:** Prepends topic from previous query 3. **"What if/How about" Scenarios:** Adds contextual parenthetical 4. **Short Vague Questions:** Expands with recent topic if pattern matches

Conversation Memory Implementation

Frontend (hooks/useChat.ts): - Maintains full message history - Passes last 10 messages to backend - Handles source mode changes mid-conversation - Preserves verification status per message

Backend (`gemini/chatService.ts`): - Receives conversation history array - Uses for both context expansion AND AI generation - Passes to Gemini 2.5 Pro for coherent responses - Logs context expansions for debugging

Memory Window: - **Short-term**: Last 10 messages (~2000 tokens) - **Long-term**: Query expansion looks at last 2 exchanges for context - **Persistent**: All messages stored client-side (no server storage)

API Integrations

1. CourtListener API v4

Purpose: Authoritative case law research **Endpoint**: `https://www.courtlistener.com/api/rest/v4/search/` **Coverage**: 1M+ California opinions (Supreme Court, Courts of Appeal) **Smart Detection**: Only activates for case law queries (contains "v.", "case", "court")

Query Optimization: - Filters for California jurisdiction - Prioritizes recent cases (2020-2025) - Returns case metadata + opinion snippets - Handles "case name" pattern matching

2. Legislative APIs (OpenStates + LegiScan)

Purpose: Current California bill tracking and full text access **Coverage**: All California bills (2023-2025 sessions) **Data Retrieved**: - Full bill text (latest version) - Bill status and legislative history - Amendments and voting records - Sponsor information

Bill Text Pipeline:

```
Query "AB 489" → Parallel API Calls (OpenStates + LegiScan)
→ Extract bill ID → Fetch full text → Decode base64 → Parse section
→ Return: { title, fullText, status, amendments }
```

3. CEB Vector Database (Upstash)

Purpose: Authoritative practice guide **RAG Architecture**: Serverless vector database with semantic search **Search Strategy**: - Query embedding generation (OpenAI) - Category-specific namespaces (5 verticals) - Cosine similarity with 0.7 confidence threshold - Metadata includes full text (10KB chunks)

Retrieval Parameters: - **Top-K**: 3-5 documents per search - **Confidence**: ~0.7 (70% similarity minimum) - **Namespace**: `ceb_${category}` (vertical isolation) - **Metadata**: Source file, page, section, CEB citation

4. Google Search Grounding (Gemini Native)

Purpose: Real-time legal updates beyond training cutoff **Activation**: Automatic for 2024-2025 queries **Search Strategy**: - Queries: "California [topic] 2025", "Governor Newsom [bill] October 2025" - Sources: Prioritizes .gov, .ca.gov, legislature websites - Coverage: Recent legislation, court decisions, regulatory changes

Grounding Metadata:

```
{  
  "webSearchQueries": ["California AI bills 2025", "SB 53 California AI  
  "groundingChunks": [  
    {  
      "web": {  
        "uri": "https://leginfo.legislature.ca.gov/faces/billNavClient.  
        "title": "SB 53 - Transparency in Frontier AI Act"  
      }  
    }  
  ]  
}
```

Deployment & Environment

Vercel Configuration

Serverless Architecture: - API routes: Node.js 18 (TypeScript) - Build: Vite + React 19 - Environment: Serverless functions (no persistent storage) - Database: Upstash Vector (serverless vector database)

Required Environment Variables:

```
# AI Services  
GEMINI_API_KEY=your_gemini_api_key_here  
ANTHROPIC_API_KEY=your_anthropic_api_key_here  
OPENAI_API_KEY=your_openai_api_key_here  
  
# CEB Vector Database  
UPSTASH_VECTOR_REST_URL=https://your-index.upstash.io  
UPSTASH_VECTOR_REST_TOKEN=your_upstash_token_here  
  
# External Legal APIs (Optional)  
COURTLISTENER_API_KEY=your_courtlistener_key_here  
OPENSTATES_API_KEY=your_openstates_key_here  
LEGISCAN_API_KEY=your_legiscan_key_here
```

Production Deployment Checklist

1. **Repository:** Connected to GitHub (automatic deploys)
2. **Environment:** All 6 variables set in Vercel dashboard
3. **Build Settings:**
 - Framework: Vite
 - Root Directory: . (project root)
 - Build Command: npm run build
 - Output Directory: dist
4. **Domain:** Custom domain or Vercel subdomain
5. **Monitoring:** Vercel logs + Sentry integration (optional)

Local Development

```
# Install dependencies  
npm install
```

```
# Start development server  
npm run dev  
  
# Build for production  
npm run build  
  
# Preview production build  
npm run preview
```

User Experience Features

1. Source Mode Selector

Location: Header section (above chat window) **Functionality:** Toggle between 3 research modes
Visual Indicators: - Active mode highlighted with primary color - Mode descriptions below selector - Real-time badge updates in conversation

2. Response Badges & Indicators

CEB Integration Badges: - **CEB VERIFIED** (Amber): Primary CEB source, no verification needed - **Hybrid Mode** (,,): Combined CEB + AI sources - **CEB Sources:** X/5 verticals covered in response

Verification Status: - **Verified:** 100% claim coverage - **Partially Verified:** 60-99% coverage with caveats - **Verification Recommended:** <60% coverage, attorney review needed

3. Source Attribution

Click-to-Expand Sources: - CEB: Cal. Prac. Guide: Family Law Â§ 3:45 (p. 127) - Legislation: AB 489 (2025) - Full bill text - Case Law: In re Marriage of Brown, 212 Cal.App.4th 967 (2013) - Web: California Courts - Official Source

Source Confidence Display: - CEB: 85-95% (semantic similarity) - Bill Text: 100% (official legislative text) - Case Law: 70-90% (CourtListener relevance) - Google Search: 60-85% (grounding metadata)

4. Conversation Memory Indicators

Context Awareness: - Follow-up questions automatically expanded - Recent topic references maintained - Conversation flow preserved across modes

Example Interaction:

User: "What is California Family Code 4320?"
Bot: [Explains spousal support factors] â€"

User: "What about the duration factor?"
Bot: [Understands this refers to Family Code 4320 duration factor] â€"

User: "How does it work for short marriages?"
Bot: [Maintains context from entire conversation] →

Technical Implementation

1. TypeScript Type Definitions (`types.ts`)

CEB Source Interface

```
export interface CEBSource extends Source {  
    isCEB: true;  
    category: 'trusts_estates' | 'family_law' | 'business_litigation' | ...  
    cebCitation: string;  
    pageNumber?: number;  
    section?: string;  
    confidence: number; // 0-1 similarity score  
}
```

Chat Message Interface

```
export interface ChatMessage {  
    id: string;  
    role: MessageRole;  
    text: string;  
    sources?: (Source | CEBSource)[];  
    verificationStatus?: VerificationStatus;  
    sourceMode?: SourceMode; // 'ceb-only' | 'ai-only' | 'hybrid'  
    isCEBBased?: boolean;  
    cebCategory?: string;  
}
```

2. CEB Search API (`api/ceb-search.ts`)

```
// Upstash Vector integration  
import { Index } from '@upstash/vector';  
  
export default async function handler(req: any, res: any) {  
    const { query, topK = 5, category } = req.body;  
  
    // Generate embedding for semantic search  
    const embedding = await generateEmbedding(query);  
  
    // Query specific CEB namespace  
    const index = new Index({  
        url: process.env.UPSTASH_VECTOR_REST_URL!,  
        token: process.env.UPSTASH_VECTOR_REST_TOKEN!  
    });  
  
    const results = await index.query({  
        vector: embedding,  
        topK,
```

```

        namespace: `ceb_${category}`,
        includeMetadata: true
    });

    // Return formatted CEB sources
    const cebSources = results.map(result => ({
        id: result.id,
        title: result.metadata.title,
        text: result.metadata.text,
        cebCitation: result.metadata.ceb_citation,
        confidence: result.metadata.confidence,
        pageNumber: result.metadata.page_number,
        category: result.metadata.category,
        // ... other metadata
    }));
}

res.status(200).json({ sources: cebSources });
}

```

3. Chat Service Orchestration (**gemini/chatService.ts**)

Query Processing Pipeline

```

async processHybrid(message: string, conversationHistory, signal) {
    // 1. Intelligent query expansion
    const expandedQuery = this.expandQueryWithContext(message, conversationHistory);

    // 2. Parallel source retrieval
    const [cebResult, aiResult] = await Promise.all([
        this.searchCEB(expandedQuery), // Vector search
        this.searchExternalSources(expandedQuery) // CourtListener + Legis
    ]);

    // 3. CEB-first context building
    const context = this.buildHybridContext(cebResult, aiResult);

    // 4. Gemini synthesis with clear source hierarchy
    const response = await this.generateWithSources(expandedQuery, context);

    // 5. Conditional verification
    const verificationStatus = this.determineVerificationNeed(cebResult, aiResult);

    return {
        text: response,
        sources: combinedSources,
        verificationStatus,
        isCEBBased: cebResult.confidence >= 0.7
    };
}

```

Context-Aware Query Expansion

```
private expandQueryWithContext(message: string, history: ConversationHistory) {
    // Pattern: "What about 460?" after "What is Penal Code 459?"
    if (/^what about|how about/i.test(message)) {
        const lastQuery = this.getLastUserQuery(history);
        const codeMatch = lastQuery.match(/(Penal Code|Family Code) \?(\d+)?/);
        if (codeMatch) {
            const numberMatch = message.match(/(\d+)/);
            if (numberMatch) {
                return `What is ${codeMatch[1]} ${numberMatch[1]}?`;
            }
        }
    }

    // Pattern: "Does it apply to...?" after legal topic
    if (/^does it|is it|can it/i.test(message)) {
        const topic = this.extractTopic(lastQuery);
        return `Regarding ${topic}, ${message}`;
    }

    return message; // Use original if no expansion needed
}
```

Performance & Cost Analysis

Response Latency Breakdown

Mode	Average P95	Notes
CEB Only	2.8s	4.1s Vector search + Gemini synthesis
AI Only	18.3s	25.7s API calls + full verification
Hybrid	12.1s	18.9s Parallel processing, conditional verification

Cost Structure

Component	Cost per Query	Monthly Estimate (100 queries/day)
Gemini 2.5 Pro	\$0.0015	\$4.50
Claude Sonnet 4.5	\$0.0030	\$9.00
CEB Embeddings	\$0.0000	\$0.00 (pre-computed)
Upstash Vector	\$0.0001	\$0.30
CourtListener	\$0.0000	\$0.00 (free tier)
Total	\$0.0046	\$13.80

Vector Database Performance

- **Query Latency:** <50ms (Upstash Vector)
- **Index Size:** 77,406 vectors, 1536 dimensions
- **Storage:** ~2.1GB total (2.8GB with metadata)
- **Throughput:** 1000+ QPS (serverless scaling)

Limitations & Compliance

Technical Limitations

1. **Verification Coverage:** 85-95% of claims can be automatically verified
2. **Multi-Turn Context:** Limited to 10 previous messages (2-3 exchanges)
3. **Source Availability:** Some recent bills may lack full text (<1 week)
4. **Jurisdiction Focus:** Optimized for California law only

Legal Compliance (California State Bar)

California Rules of Professional Conduct Compliance:

Rule 1.1 - Competence

- **Implemented:** Multi-source verification ensures responses are based on authoritative sources
- **Limitation:** Users must still exercise professional judgment

Rule 1.6 - Confidentiality

- **Implemented:** No persistent storage of user queries
- **Warning:** Queries transmitted to third-party AI providers (Google, Anthropic)
- **Recommendation:** Anonymize client data before input

Rule 8.4 - Misconduct

- **Implemented:** Clear disclaimers that this is not legal advice
- **Implemented:** Verification warnings for unverified claims
- **Implemented:** Source citations for all legal information

Usage Guidelines for Attorneys

MANDATORY: 1. **Anonymize all client data** before entering queries 2. **Verify all critical information** against primary sources 3. **Consult with qualified counsel** before relying on responses 4. **Review California State Bar** guidelines for AI use 5. **Check local court rules** for AI disclosure requirements

RECOMMENDED: 1. Use **CEB Only** mode for practice guidance 2. Use **Hybrid** mode for comprehensive research 3. Always click source links to verify original documents 4. Document AI use in client files 5. Train staff on AI limitations and ethical considerations

Error Handling & Fallbacks

Graceful Degradation: - API failures trigger fallback to alternative sources - Missing bill text uses Google Search grounding - Verification failures show "Verification Recommended" badge - Network errors display helpful error messages

Monitoring: - All API failures logged to Vercel dashboard - Response times tracked per component - Source availability monitored (CourtListener, Upstash) - Error rates per source type recorded

Appendix: System Specifications

Model Specifications

Model	Provider	Role	Context Window	Temperature
Gemini 2.5 Pro	Google	Generator	1M tokens	0.2
Claude Sonnet 4.5	Anthropic	Verifier	200K tokens	0.0
text-embedding-3-small	OpenAI	Embeddings	N/A	N/A

API Rate Limits

Service	Limit	Implementation
Upstash Vector	1000 QPS	Automatic retry + exponential backoff
CourtListener	100/hour	Smart caching (5 minutes)
OpenStates	1000/hour	Query deduplication
Google Search	60/minute	Built into Gemini API
Claude API	100/minute	Request queuing

Data Processing Pipeline

Hardware Requirements (Processing): - **CPU**: 8+ cores recommended - **RAM**: 16GB+ (M4 Max 128GB optimal) - **GPU**: Optional for embedding generation (accelerates OpenAI calls)

Processing Time (Initial Load):

Step	Duration	Parallelizable
PDF Processing	2-3 hours	Yes (per vertical)
Embedding Generation	1-2 hours	Yes (batch processing)
Vector Upload	28 minutes	Yes (batch upsert)
Total	3-5 hours	High

Cost Analysis (Ongoing)

Monthly Operational Costs (100 queries/day): - **Gemini 2.5 Pro**: \$4.50 (primary generation) - **Claude Sonnet 4.5**: \$9.00 (verification) - **Upstash Vector**: \$0.30 (search storage) - **OpenAI Embeddings**: \$0.00 (pre-computed) - **Vercel Hosting**: \$20.00 (serverless) - **Total: \$33.80/month**

Cost per Query: \$0.0046 (enterprise-grade legal research)

Last Updated: November 3, 2025

Version: 2.1 (CEB Integration Complete)

Author: AI Development Team

License: MIT